

Assignment 3

- Due Nov 30, 2021 by 11:05pm
- Points 190
- Submitting a file upload
- File Types zip
- Attempts 1
- Allowed Attempts 3
- Available Nov 1, 2021 at 2:30pm - Nov 30, 2021 at 11:05pm

This assignment was locked Nov 30, 2021 at 11:05pm.

Homework Assignment 3

EU Cities Temperatures, German Credit, Google Playstore Apps

Worth 190 points.

Posted Monday, November 1

Due Tuesday, November 30 at 11 PM in Canvas

NO LATE SUBMISSIONS ACCEPTED

Everyone will work on the assignment in pairs. **You DO NOT have the option of working alone.**

Make sure you abide by the [DCS Academic Integrity Policy for Programming Assignments](https://www.cs.rutgers.edu/academics/undergraduate/academic-integrity-policy/programming-assignments). 
<https://www.cs.rutgers.edu/academics/undergraduate/academic-integrity-policy/programming-assignments>




For the problems in this homework, write your code in Jupyter notebooks: **eucities_temps.ipynb**, **german_credit.ipynb**, and **google_apps.ipynb** respectively.

In each notebook, write appropriate markdowns for each task (just copy the task description from the assignment) so we know which task you are doing in the cells following the markdown.

Zip all of these notebooks into a single file named **hw3.zip** and submit this to Canvas. Do not include any input or output files, we only need your notebooks.
ONLY ONE SUBMISSION PER GROUP, PLEASE!

You are allowed up to 3 submissions, **only the last submission will be graded.**

Problem 1: EU Cities Temperatures Dataset (55 points)

Given a CSV data file as represented by the sample file [EuCitiesTemperatures.csv](https://rutgers.instructure.com/courses/133454/files/20252138?wrap=1) (https://rutgers.instructure.com/courses/133454/files/20252138/download?download_frd=1)  https://docreader.readspeaker.com/docreader/?cid=8909&lang=en_us&url=https%3A%2F%2Finst-fs-iad-prod.inscloudgate.net%2Ffiles%2F0380fdf4-c845-4ed9-a773-c9707b2fa26e%2FEuCitiesTemperatures.csv%3Ftoken%3DeyJ0eXAiOiJKV1QiLCJhbGciOiJIUzUxMiJ9.eyJpYXQiOiE3MjYyNzQ0NDQsInVzZXJfaWQiOiE5bGwslInJlE7ko8g_j5Oz8SdZM0MqH6DKKK5FIHhSmLcm07I4lmF2Lqai24ijoAGXcuYDEkpHqsrCDw. (213 records), load it into a Pandas DataFrame and perform the following tasks on it.

Important: Your code should be applicable to any extension of this sample, so make sure you don't hardcode anything that applies only to the values in this dataset.

Preprocessing/Analysis (28 pts)

1. [9 pts] Fill in the missing latitude and longitude values by calculating the average for that country. Round the average to 2 decimal places.
2. [9 pts] Find out the subset of cities that lie between latitudes 40 to 60 (both inclusive) and longitudes 15 to 30 (both inclusive). Find out which country have the maximum number of cities in this geographical band. (More than one country could have the maximum number of values.)
3. [10 pts] Fill in the missing temperature values by the average temperature value of the similar region type. A region type would be a combination of whether it is in EU (yes/no) and whether it has a coastline (yes/no).



For example, if we have a missing temperature value for Bergen, Norway, which is not in the EU but lies on the coast, we will fill it with the average temperature of cities with EU='no' and coastline='yes')

Visualization (27 pts)

For all plots, make sure to label the axes, and set appropriate tick labels.

- [6 pts] Plot a bar chart for the number of cities belonging to each of the regions described in Preprocessing/Analysis #3 above.
- [7 pts] Plot a scatter plot of latitude (y-axis) v/s longitude (x-axis) values to get a map-like visual of the cities under consideration. All the cities in the same country should have the same color.
- [6 pts] The population column contains values unique to each country. So two cities of the same country will show the same population value. Plot a histogram of the number of countries belonging to each population group: split the population values into 5 bins (groups).
- [8 pts] Plot subplots (2, 2), with proper titles, one each for the region types described in Preprocessing/Analysis #3 above.

Each subplot should be a scatter plot of Latitude (y-axis) vs. City (x-axis), where the color of the plot points should be based on the temperature values: 'red' for temperatures above 10, 'blue' for temperatures below 6 and 'orange' for temperatures between 6 and 10 (both inclusive). For each subplot, set xticks to an array of numbers from 0 to n-1 (both inclusive), where n is the total number of cities in each region type. This represents each city as a number between 0 and n-1.

Problem 2: German Credit Dataset (72 points)

Given a CSV data file as represented by the sample file [GermanCredit.csv](https://rutgers.instructure.com/courses/133454/files/20252574?wrap=1) (https://rutgers.instructure.com/courses/133454/files/20252574/download?download_frd=1), https://docreader.readspeaker.com/docreader/?cid=8909&lang=en_us&url=https%3A%2F%2Finst-fs-iad-prod.inscloudgate.net%2Ffiles%2F56210f9e-9b9b-49ea-bdb5-b1dab6587432%2FGermanCredit.csv%3Ftoken%3DeyJ0eXAiOiJKV1QiLCJhbGciOiJIUzUxMiJ9.eyJpYXQiOiE3MjYyNzY2MDEsInVzZXJfaWQiOm51bGwsInJlc291cm-qg6qlvuVVTB7bCALOZJ-V80TgfOu9iNegw4GczpH7ZpXU-E3EC05q5z5q7uGpzx2vCM6Q) (1000 records), load it into a Pandas DataFrame, and perform the following tasks on it.

Important: Your code should be applicable to any extension of this sample, so make sure you don't hardcode anything that applies only to the values in this sample.

Preprocessing (31 pts)

- [8 pts] Drop the 3 columns that contribute the least to the dataset. These would be the columns with the highest number of non-zero 'none' values. Break ties by going left to right in columns. (Your code should be generalizable to drop n columns, but for the rest of the analysis, you can call your code for n=3.)
- [4 pts] Certain values in some of the columns contain unnecessary apostrophes ('). Remove the apostrophes.
- [5 pts] The checking_status column has values in 4 categories: 'no checking', '<0', '0<=X<200', and '>=200'. Change these to 'No Checking', 'Low', 'Medium', and 'High' respectively.
- [5 pts] The savings_status column has values in 4 categories: 'no known savings', '<100', '100<=X<500', '500<=X<1000', and '>=1000'. Change these to 'No Savings', 'Low', 'Medium', 'High', and 'High' respectively. (Yes, the last two are both 'High').
- [4 pts] Change class column values from 'good' to '1' and 'bad' to '0'
- [5 pts] Change the employment column value 'unemployed' to 'Unemployed', and for the others, change to 'Amateur', 'Professional', 'Experienced' and 'Expert', depending on year range.

Analysis (17 pts)

For the following tasks, do preprocessing or changing of data types in the data frame as required.

- [5 pts] Often we need to find correlations between categorical attributes, i.e. attributes that have values that fall in one of several categories, such as "yes"/"no" for attr1, or "low", "medium", "high" for attr2.

One such correlation is to find counts in combinations of categorical values across attributes, as in how many instances are "yes" for attr1 and "low" for attr2. A good way to find such counts is to use the Pandas [crosstab](https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.crosstab.html) function. Do this for the following two counts.

- [3 pts] Get the count of each category of foreign workers (yes and no) for each class of credit (good and bad).
- [2 pts] Similarly, get the count of each category of employment for each category of saving_status.



- ```
Most common checking status: ...
Most common savings status: ...
```



- a. Rating (gives top 3 most highly rated applications in each category)
- b. Installs (gives top 3 most installed applications in each category)
- c. Reviews (gives top 3 most reviewed applications in each category)

You don't need to do anything explicit to break ties.

Each of the returned dataframes have Category and App for the first two columns, and one of Rating (for a.), Installs (for b.), and Reviews (for c.) as the third column.

3. [4 pts] Find the average, maximum and minimum price of the paid applications.

## Visualization (16 pts)

1. [9 pts] In the genre column, break the string of genres into a list. For example, 'Art & Design; Creativity' should be ['Art & Design', Creativity].

Count the number of applications per genre and display it using a pie chart.

**Hint:** Read about `DataFrame.explode()`

2. [7 pts] Display a box plot of ratings for "Business" and "Education" categories. The boxplots should be in the same plot.
- 

